

# CS145 Homework 1

**Important Note:** HW1 is due on **11:59 PM PT, Oct 19 (Monday, Week 3)**. Please submit through GradeScope (you will receive an invite to Gradescope for CS145 Fall 2020.).

## Print Out Your Name and UID

Name: Kevin Li, UID: 405200619

## Before You Start

You need to first create HW1 conda environment by the given `cs145hw1.yml` file, which provides the name and necessary packages for this tasks. If you have `conda` properly installed, you may create, activate or deactivate by the following commands:

```
conda env create -f cs145hw1.yml
conda activate hw1
conda deactivate
```

OR

```
conda env create --name NAMEOFOURCHOICE -f cs145hw1.yml
conda activate NAMEOFOURCHOICE
conda deactivate
```

To view the list of your environments, use the following command:

```
conda env list
```

More useful information about managing environments can be found [here](https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html) (<https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html>).

You may also quickly review the usage of basic Python and Numpy package, if needed in coding for matrix operations.

In this notebook, you must not delete any code cells in this notebook. If you change any code outside the blocks that you are allowed to edit (between `START/END YOUR CODE HERE`), you need to highlight these changes. You may add some additional cells to help explain your results and observations.

```
In [19]: import numpy as np
import pandas as pd
import sys
import random as rd
import matplotlib.pyplot as plt
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:  
%reload\_ext autoreload

If you can successfully run the code above, there will be no problem for environment setting.

## 1. Linear regression

This workbook will walk you through a linear regression example.

```
In [20]: from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
# As a sanity check, we print out the size of the training data (1000, 100) and test data (100, 100)
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape: ', lm.train_y.shape)
```

Training data shape: (1000, 100)  
Training labels shape: (1000,)

### 1.1 Closed form solution

In this section, complete the `getBeta` function in `linear_regression.py` which use the closed form solution of  $\hat{\beta}$ .

Train your model by using `lm.train('0')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```
In [21]: from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#####
# STRART YOUR CODE HERE #
#####

beta = lm.train('0')
y_train_hat = lm.predict(lm.train_x, beta)
y_test_hat = lm.predict(lm.test_x, beta)

training_error = lm.compute_mse(y_train_hat, lm.train_y)
testing_error = lm.compute_mse(y_test_hat, lm.test_y)

#####
# END YOUR CODE HERE #
#####
print('Training error is: ', training_error)
print('Testing error is: ', testing_error)
```

```
Learning Algorithm Type: 0
Training error is: 0.08693886675396781
Testing error is: 0.11017540281675806
```

## 1.2 Batch gradient descent

In this section, complete the `getBetaBatchGradient` function in `linear_regression.py` which compute the gradient of the objective fuction.

Train you model by using `lm.train('1')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```

In [22]: lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#=====#
# STRART YOUR CODE HERE #
#=====#

beta = lm.train('1')
y_train_hat = lm.predict(lm.train_x, beta)
y_test_hat = lm.predict(lm.test_x, beta)

training_error = lm.compute_mse(y_train_hat, lm.train_y)
testing_error = lm.compute_mse(y_test_hat, lm.test_y)

#=====#
# END YOUR CODE HERE #
#=====#
print('Training accuracy is: ', training_error)
print('Testing accuracy is: ', testing_error)

```

```

Learning Algorithm Type: 1
Training accuracy is: 0.08693902047792586
Testing accuracy is: 0.11018845441523527

```

### 1.3 Stochastic gadient descent

In this section, complete the `getBetaStochasticGradient` function in `linear_regression.py`, which use an estimated gradient of the objective function.

Train you model by using `lm.train('2')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```

In [23]: lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#####
# STRART YOUR CODE HERE #
#####

lm.lr = 0.0005
beta = lm.train('2')
y_train_hat = lm.predict(lm.train_x, beta)
y_test_hat = lm.predict(lm.test_x, beta)

training_error = lm.compute_mse(y_train_hat, lm.train_y)
testing_error = lm.compute_mse(y_test_hat, lm.test_y)

#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_error)
print('Testing accuracy is: ', testing_error)

```

Learning Algorithm Type: 2  
 Training accuracy is: 0.09531517836589815  
 Testing accuracy is: 0.11959236638904801

## Questions:

1. Compare the MSE on the testing dataset for each version. Are they the same? Why or why not?
2. Apply z-score normalization for eachh featurre and comment whether or not it affect the three algorithm.
3. Ridge regression is adding an L2 regularization term to the original objective function of mean squared error. The objective function become following:

$$J(\beta) = \frac{1}{2n} \sum_i (x_i^T \beta - y_i)^2 + \frac{\lambda}{2n} \sum_j \beta_j^2,$$

where  $\lambda \geq 0$ , which is a hyper parameter that controls the trade off. Take the derivative of this provided objective function and derive the closed form solution for  $\beta$ .

## Your answer here:

```
In [24]: # Running normalized versions
lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
print(f'Feature average std: {lm.train_x.describe().loc["std"].mean()}\n')
lm.normalize()

def run_alg(a):
    beta = lm.train(a)
    training_error = lm.compute_mse(lm.predict(lm.train_x, beta), lm.train_y)
    testing_error = lm.compute_mse(lm.predict(lm.test_x, beta), lm.test_y)
    print('Training error is: ', training_error)
    print('Testing error is: ', testing_error, '\n')

for a in ('0', '1', '2'):
    run_alg(a)
```

Feature average std: 2.3522292509877913

Learning Algorithm Type: 0  
 Training error is: 0.08693886675396784  
 Testing error is: 0.11017540281675804

Learning Algorithm Type: 1  
 Training error is: 0.10018132194244624  
 Testing error is: 0.13900607116647945

Learning Algorithm Type: 2  
 Training error is: 0.10992817207625191  
 Testing error is: 0.13117156993305695

1. The MSE is slightly different but close for each version. They're different because they each take different paths to optimizing beta, but they all end up in a similar place because (it seems like) this problem is probably convex i.e. only has one optimum solution. To be precise, the first closed form solution should actually get the beta with the minimum error while BGD and SGD simply run for a certain number of iterations, then stop wherever they are. Since they involve choosing a random beta and SGD randomly samples from the training dataset, they will perform slightly randomly (although asymptotically they should still converge on the optimal beta)
2. Z-scoring the features has no effect on the closed form solution but does slightly affect the gradient descent algorithms. This is because linear regression simply assigns a weight ( $\beta_i$ ) to each feature; if the feature's size is scaled, then the respective weight will be scaled appropriately to compensate. The "zeroing" of the features is also accounted for due to the addition of the bias term (which we see is added via the addAllOneColumn function). For the gradient descent algorithms, since they actually use the feature values to calculate their gradients and are affected by the feature values, we see mild effects on the model. One reason this happens is because Z-scoring (normalizing) is done to make each feature "just as important" as each other to avoid quickly converging on a solution that only uses some features. For example, if the standard deviation of  $x_1$  is massive and the standard deviation of  $x_2$  is tiny and we had the true solution be  $y = x_1 + x_2$ , then we would be able to converge on a "pretty good" solution by simply guessing  $y = x_1$  i.e.  $b_1 = 1, b_2 = 0$ . Note that this does not necessarily mean normalization improves performance! In fact in our case, it seems

to hurt performance (although it would be nice to do something more like k-fold validation to investigate more thoroughly). It also helps with avoiding overflow which we saw affect SGD (previously without normalization, we needed to reduce the learning rate to avoid overflow).

3. The solution ends up being  $\beta = (X^T X + \lambda I)^{-1} X^T y$

$$J(\beta) = \frac{1}{2n} \sum_i (x_i^T \beta - y_i)^2 + \frac{\lambda}{2n} \sum_j \beta_j^2$$

$$J(\beta) = \frac{1}{2n} (X\beta - Y)^T (X\beta - Y) + \frac{\lambda}{2n} \beta^T \beta$$

We know what the left side (original loss) turns into since it's the same as the original loss function  $J(\beta)$ , so we get the following

$$\frac{\partial J}{\partial \beta} = (X^T X \beta - X^T y)/n + \frac{\partial J}{\partial \beta} \left( \frac{\lambda}{2n} \beta^T \beta \right)$$

Taking the derviative of the right side, we get

$$\frac{\partial J}{\partial \beta} = (X^T X \beta - X^T y)/n + \frac{\lambda}{n} \beta$$

We set the derviative equal to zero and do some rearranging and get

$$0 = \frac{1}{n} (X^T X \beta - X^T y + \lambda \beta)$$

$$X^T y = X^T X \beta + \lambda \beta$$

$$X^T y = \beta (X^T X + \lambda I)$$

$$(X^T X + \lambda I)^{-1} X^T y = \beta$$

## 2. Logistic regression

This workbook will walk you through a logistic regression example.

```
In [25]: from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-test.csv')
# As a sanity check, we print out the size of the training data (1000, 5) and test data (1000, 5)
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape:', lm.train_y.shape)
```

Training data shape: (1000, 5)

Training labels shape: (1000,)

### 2.1 Batch gradiend descent

In this section, complete the `getBeta_BatchGradient` in `logistic_regression.py`, which compute the gradient of the log likelihoood function.

Complete the `compute_avglogL` function in `logistic_regression.py` for sanity check.

Train you model by using `lm.train('0')` function.

And print the training and testing accuracy using `lm.predict` and `lm.compute_accuracy` given.

```
In [26]: lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-test.csv')
training_accuracy= 0
testing_accuracy= 0
#####
# STRART YOUR CODE HERE #
#####

lm.normalize()

beta = lm.train('0')
y_train_hat = lm.predict(lm.train_x, beta)
y_test_hat = lm.predict(lm.test_x, beta)

training_accuracy = lm.compute_accuracy(y_train_hat, lm.train_y)
testing_accuracy = lm.compute_accuracy(y_test_hat, lm.test_y)

#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_accuracy)
print('Testing accuracy is: ', testing_accuracy)
```

average logL for iteration 0: -0.4893882425713696  
average logL for iteration 1000: -0.460100375350853  
average logL for iteration 2000: -0.460100375350853  
average logL for iteration 3000: -0.460100375350853  
average logL for iteration 4000: -0.460100375350853  
average logL for iteration 5000: -0.460100375350853  
average logL for iteration 6000: -0.460100375350853  
average logL for iteration 7000: -0.460100375350853  
average logL for iteration 8000: -0.460100375350853  
average logL for iteration 9000: -0.460100375350853  
Training avgLogL: -0.460100375350853  
Training accuracy is: 0.797  
Testing accuracy is: 0.7534791252485089

## 2.2 Newton Raphhson

In this section, complete the `getBeta_Newton` in `logistic_regression.py`, which make use of both first and second derivative.

Train you model by using `lm.train('1')` function.

Print the training and testing accuracy using `lm.predict` and `lm.compute_accuracy` given.



```

In [27]: lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-t
training_accuracy= 0
testing_accuracy= 0
#####
# STRART YOUR CODE HERE #
#####

lm.normalize()

beta = lm.train('1')
y_train_hat = lm.predict(lm.train_x, beta)
y_test_hat = lm.predict(lm.test_x, beta)

training_accuracy = lm.compute_accuracy(y_train_hat, lm.train_y)
testing_accuracy = lm.compute_accuracy(y_test_hat, lm.test_y)

#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_accuracy)
print('Testing accuracy is: ', testing_accuracy)

```

```

average logL for iteration 0: -0.4905626329298569
average logL for iteration 500: -0.460100375350853
average logL for iteration 1000: -0.460100375350853
average logL for iteration 1500: -0.460100375350853
average logL for iteration 2000: -0.460100375350853
average logL for iteration 2500: -0.460100375350853
average logL for iteration 3000: -0.460100375350853
average logL for iteration 3500: -0.460100375350853
average logL for iteration 4000: -0.460100375350853
average logL for iteration 4500: -0.460100375350853
average logL for iteration 5000: -0.460100375350853
average logL for iteration 5500: -0.460100375350853
average logL for iteration 6000: -0.460100375350853
average logL for iteration 6500: -0.460100375350853
average logL for iteration 7000: -0.460100375350853
average logL for iteration 7500: -0.460100375350853
average logL for iteration 8000: -0.460100375350853
average logL for iteration 8500: -0.460100375350853
average logL for iteration 9000: -0.460100375350853
average logL for iteration 9500: -0.460100375350853
Training avgLogL: -0.460100375350853
Training accuracy is: 0.797
Testing accuracy is: 0.7534791252485089

```

## Questions:

1. Compare the accuracy on the testing dataset for each version. Are they the same? Why or why not?
2. Regularization. Similar to linear regression, an regularization term could be added to logistic regression. The objective function becomes following:

$$J(\beta) = -\frac{1}{n} \sum_i (y_i x_i^T \beta - \log(1 + \exp\{x_i^T \beta\})) + \lambda \sum_j \beta_j^2,$$

where  $\lambda \geq 0$ , which is a hyper parameter that controls the trade off. Take the derivative  $\frac{\partial J(\beta)}{\partial \beta_j}$  of this provided objective function and provide the batch gradient descent update.

### Your answer here:

1. They are the same since the both converged on the same (probably very close to optimal) solution. Newton-Raphson is simply a different optimizer and can therefore have different properties (i.e. speed, ability to find global optimum, precision, etc) but if they both find the global optimum, they will obviously have the same results
2. The batch gradient update ends up being  $\beta_j^{t+1} = B_j^t + \eta (-2\lambda\beta_j + \frac{1}{n} \sum_i x_{ij}(y_i - p_i(\beta)))$

$$J(\beta) = -\frac{1}{n} \sum_i (y_i x_i^T \beta - \log(1 + e^{x_i^T \beta})) + \lambda \sum_j \beta_j^2$$

Like before, we know the derivative of the left part of the equation since it's the regular log likelihood loss function and we covered the derivative in class, so we gradient

$$\begin{aligned} \frac{\partial J}{\partial \beta_j} &= -\frac{1}{n} \sum_{i=1}^n x_{ij}(y_i - p_i(\beta)) + \frac{\partial}{\partial \beta_j} \lambda \sum_j \beta_j^2 \\ \frac{\partial J}{\partial \beta_j} &= -\frac{1}{n} \sum_{i=1}^n x_{ij}(y_i - p_i(\beta)) + 2\lambda\beta_j \end{aligned}$$

Now that we know the derivative, gradient descent becomes as simple as

$$\begin{aligned} \beta_j^{t+1} &= B_j^t - \eta \frac{\partial J}{\partial \beta_j} \\ \beta_j^{t+1} &= B_j^t - \eta \left( -\frac{1}{n} \sum_i x_{ij}(y_i - p_i(\beta)) + 2\lambda\beta_j \right) \\ \beta_j^{t+1} &= B_j^t + \eta \left( -2\lambda\beta_j + \frac{1}{n} \sum_i x_{ij}(y_i - p_i(\beta)) \right) \end{aligned}$$

## 2.3 Visualize the decision boundary on a toy dataset

In this subsection, you will use the same implementation for another small dataset with each datapoint  $x$  with only two features ( $x_1, x_2$ ) to visualize the decision boundary of logistic regression model.

```
In [29]: from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression(verbose = False)
lm.load_data('./data/logistic-regression-toy.csv', './data/logistic-regression-toy')
# As a sanity check, we print out the size of the training data (99,2) and training labels
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape:', lm.train_y.shape)
```

```
Training data shape: (99, 2)
Training labels shape: (99,)
```

In the following block, you can apply the same implementation of logistic regression model (either in 2.1 or 2.2) to the toy dataset. Print out the  $\hat{\beta}$  after training and accuracy on the train set.

```
In [30]: training_accuracy= 0
#=====#
#  START YOUR CODE HERE  #
#=====#

lm.normalize()
beta = lm.train('1')
y_train_hat = lm.predict(lm.train_x, beta)
y_test_hat = lm.predict(lm.test_x, beta)

training_accuracy = lm.compute_accuracy(y_train_hat, lm.train_y)
testing_accuracy = lm.compute_accuracy(y_test_hat, lm.test_y)

print(f'Beta: {beta}')
```

```
Training avgLogL: -0.329147431295712
Beta: [-0.04717577  1.46005896  2.06586134]
Training accuracy is: 0.8888888888888888
```

Next, we try to plot the decision boundary of your learned logistic regression classifier. Generally, a decision boundary is the region of a space in which the output label of a classifier is ambiguous. That is, in the given toy data, given a datapoint  $x = (x_1, x_2)$  on the decision boundary, the logistic regression classifier cannot decide whether  $y = 0$  or  $y = 1$ .

## Question

Is the decision boundary for logistic regression linear? Why or why not?

## Your answer here:

Yes.

We know the decision boundary for logistic regression is when  $\sigma(x^T \beta) = 0.5$ . Then, we know that

$$\begin{aligned}\sigma(x^T \beta) &= 0.5 \\ \frac{1}{1 + e^{-x^T \beta}} &= 0.5 \\ 1 &= e^{-x^T \beta} \\ 0 &= -x^T \beta\end{aligned}$$

This is the equation for a hyperplane - the decision boundary is linear

Draw the decision boundary in the following cell. Note that the code to plot the raw data points are given. You may need `plt.plot` function (see [here](https://matplotlib.org/tutorials/introductory/pyplot.html) (<https://matplotlib.org/tutorials/introductory/pyplot.html>)).


```
In [35]: # scatter plot the raw data
df = pd.concat([lm.train_x, lm.train_y], axis=1)
groups = df.groupby("y")
for name, group in groups:
    plt.plot(group["x1"], group["x2"], marker="o", linestyle="", label=name)

# plot the decision boundary on top of the scattered points
#=====#
# START YOUR CODE HERE #
#=====#

x = np.array([-2, 2])
y = (beta[0] + beta[1] * x) / -beta[2]

plt.plot(x, y)

#=====#
# END YOUR CODE HERE #
#=====#
plt.show()
```



## End of Homework 1 :)

After you've finished the homework, please print out the entire `ipynb` notebook and two `py` files into one PDF file. Make sure you include the output of code cells and answers for questions. Prepare submit it to GradeScope.

```

1 import pandas as pd
2 import numpy as np
3 import sys
4 import random as rd
5
6 #insert an all-one column as the first column
7 def addAllOneColumn(matrix):
8     n = matrix.shape[0] #total of data points
9     p = matrix.shape[1] #total number of attributes
10
11     newMatrix = np.zeros((n,p+1))
12     newMatrix[:,1:] = matrix
13     newMatrix[:,0] = np.ones(n)
14
15     return newMatrix
16
17 # Reads the data from CSV files, converts it into Dataframe and returns x and y
18 # dataframes
19 def getDataframe(filePath):
20     dataframe = pd.read_csv(filePath)
21     y = dataframe['y']
22     x = dataframe.drop('y', axis=1)
23     return x, y
24
25 # train_x and train_y are numpy arrays
26 # function returns value of beta calculated using (0) the formula  $\beta = (X^T X)^{-1} (X^T Y)$ 
27 def getBeta(train_x, train_y):
28     n = train_x.shape[0] #total of data points
29     p = train_x.shape[1] #total number of attributes
30
31     beta = np.zeros(p)
32     #=====#
33     # STRART YOUR CODE HERE #
34     #=====#
35     beta = np.linalg.inv(train_x.T @ train_x) @ (train_x.T @ train_y)
36     #=====#
37     # END YOUR CODE HERE #
38     #=====#
39     return beta
40
41 # train_x and train_y are numpy arrays
42 # lr (learning rate) is a scalar
43 # function returns value of beta calculated using (1) batch gradient descent
44 def getBetaBatchGradient(train_x, train_y, lr, num_iter):
45     beta = np.random.rand(train_x.shape[1])
46
47     n = train_x.shape[0] #total of data points
48     p = train_x.shape[1] #total number of attributes
49
50     beta = np.random.rand(p)
51     #update beta iteratively
52     for iter in range(0, num_iter):
53         deriv = np.zeros(p)
54         for i in range(n):
55             #=====#
56             # STRART YOUR CODE HERE #
57             #=====#
58             deriv += train_x[i] * (train_x[i] @ beta - train_y[i])

```

```

59         #=====#
60         #   END YOUR CODE HERE   #
61         #=====#
62         deriv = deriv / n
63         beta = beta - deriv.dot(lr)
64         return beta
65
66 # train_x and train_y are numpy arrays
67 # lr (learning rate) is a scalar
68 # function returns value of beta calculated using (2) stochastic gradient descent
69 def getBetaStochasticGradient(train_x, train_y, lr):
70     n = train_x.shape[0] #total of data points
71     p = train_x.shape[1] #total number of attributes
72
73     beta = np.random.rand(p)
74
75     epoch = 100
76     for iter in range(epoch):
77         indices = list(range(n))
78         rd.shuffle(indices)
79         for i in range(n):
80             idx = indices[i]
81             #=====#
82             #   STRART YOUR CODE HERE   #
83             #=====#
84             beta += lr * train_x[idx] * (train_y[idx] - train_x[idx] @ beta)
85             #=====#
86             #   END YOUR CODE HERE   #
87             #=====#
88         return beta
89
90
91 # Linear Regression implementation
92 class LinearRegression(object):
93     # Initializes by reading data, setting hyper-parameters, and forming linear model
94     # Forms a linear model (learns the parameter) according to type of beta (0 -
95     # closed form, 1 - batch gradient, 2 - stochastic gradient)
96     # Performs z-score normalization if z_score is 1
97     def __init__(self, lr=0.005, num_iter=1000):
98         self.lr = lr
99         self.num_iter = num_iter
100        self.train_x = pd.DataFrame()
101        self.train_y = pd.DataFrame()
102        self.test_x = pd.DataFrame()
103        self.test_y = pd.DataFrame()
104        self.algType = 0
105        self.isNormalized = 0
106
107    def load_data(self, train_file, test_file):
108        self.train_x, self.train_y = getDataframe(train_file)
109        self.test_x, self.test_y = getDataframe(test_file)
110
111    def normalize(self):
112        # Applies z-score normalization to the dataframe and returns a normalized
113        # dataframe
114        self.isNormalized = 1
115        means = self.train_x.mean(0)
116        std = self.train_x.std(0)
117        self.train_x = (self.train_x - means).div(std)
118        self.test_x = (self.test_x - means).div(std)

```

```
117
118     # Gets the beta according to input
119     def train(self, algType):
120         self.algType = algType
121         newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-one column
as the first column
122         print('Learning Algorithm Type: ', algType)
123         if(algType == '0'):
124             beta = getBeta(newTrain_x, self.train_y.values)
125             #print('Beta: ', beta)
126
127         elif(algType == '1'):
128             beta = getBetaBatchGradient(newTrain_x, self.train_y.values, self.lr,
self.num_iter)
129             #print('Beta: ', beta)
130         elif(algType == '2'):
131             beta = getBetaStochasticGradient(newTrain_x, self.train_y.values,
self.lr)
132             #print('Beta: ', beta)
133         else:
134             print('Incorrect beta_type! Usage: 0 - closed form solution, 1 - batch
gradient descent, 2 - stochastic gradient descent')
135
136
137         return beta
138
139     # Predicts the y values on given data and learned beta
140     def predict(self,x, beta):
141         newTest_x = addAllOneColumn(x)
142         self.predicted_y = newTest_x.dot(beta)
143         return self.predicted_y
144
145
146     # predicted_y and y are the predicted and actual y values respectively as numpy
arrays
147     # function returns the mean squared error (MSE) value for the test dataset
148     def compute_mse(self,predicted_y, y):
149         mse = np.sum((predicted_y - y)**2)/predicted_y.shape[0]
150         return mse
151
152
153
```



```

1  # -*- coding: utf-8 -*-
2
3  import pandas as pd
4  import numpy as np
5  import sys
6  import random as rd
7
8  #insert an all-one column as the first column
9  def addAllOneColumn(matrix):
10     n = matrix.shape[0] #total of data points
11     p = matrix.shape[1] #total number of attributes
12
13     newMatrix = np.zeros((n,p+1))
14     newMatrix[:,0] = np.ones(n)
15     newMatrix[:,1:] = matrix
16
17
18     return newMatrix
19
20 # Reads the data from CSV files, converts it into Dataframe and returns x and y
   dataframes
21 def getDataframe(filePath):
22     dataframe = pd.read_csv(filePath)
23     y = dataframe['y']
24     x = dataframe.drop('y', axis=1)
25     return x, y
26
27 # sigmoid function
28 def sigmoid(z):
29     return 1 / (1 + np.exp(-z))
30
31 # compute average logL
32 def compute_avglogL(X,y,beta):
33     eps = 1e-50
34     n = y.shape[0]
35     avglogL = 0
36     #=====#
37     # STRART YOUR CODE HERE #
38     #=====#
39     avglogL = np.sum(y * (X @ beta) - np.log(1 + np.exp(X @ beta))) / n
40     #=====#
41     # END YOUR CODE HERE #
42     #=====#
43     return avglogL
44
45
46 # train_x and train_y are numpy arrays
47 # lr (learning rate) is a scalar
48 # function returns value of beta calculated using (0) batch gradient descent
49 def getBeta_BatchGradient(train_x, train_y, lr, num_iter, verbose):
50     beta = np.random.rand(train_x.shape[1])
51
52     n = train_x.shape[0] #total of data points
53     p = train_x.shape[1] #total number of attributes
54
55
56     beta = np.random.rand(p)
57     #update beta iteratively
58     for iter in range(0, num_iter):
59         #=====#

```

```

60     # STRART YOUR CODE HERE  #
61     #=====#
62
63     beta += lr * (train_x.T @ (train_y - sigmoid(train_x @ beta)))
64
65     #=====#
66     #   END YOUR CODE HERE   #
67     #=====#
68     if(verbose == True and iter % 1000 == 0):
69         avgLogL = compute_avglogL(train_x, train_y, beta)
70         print(f'average logL for iteration {iter}: {avgLogL} \t')
71     return beta
72
73 # train_x and train_y are numpy arrays
74 # function returns value of beta calculated using (1) Newton-Raphson method
75 def getBeta_Newton(train_x, train_y, num_iter, verbose):
76     n = train_x.shape[0] #total of data points
77     p = train_x.shape[1] #total number of attributes
78
79     beta = np.random.rand(p)
80     for iter in range(0, num_iter):
81         #=====#
82         # STRART YOUR CODE HERE  #
83         #=====#
84
85         y_pred = sigmoid(train_x @ beta)
86         grad = train_x.T @ (train_y - sigmoid(train_x @ beta))
87
88         neg_hessian = train_x.T @ np.diag(y_pred * (1 - y_pred)) @ train_x
89         beta += np.linalg.inv(neg_hessian) @ grad
90
91         #=====#
92         #   END YOUR CODE HERE   #
93         #=====#
94         if(verbose == True and iter % 500 == 0):
95             avgLogL = compute_avglogL(train_x, train_y, beta)
96             print(f'average logL for iteration {iter}: {avgLogL} \t')
97     return beta
98
99
100
101 # Logistic Regression implementation
102 class LogisticRegression(object):
103     # Initializes by reading data, setting hyper-parameters
104     # Learns the parameter using (0) Batch gradient (1) Newton-Raphson
105     # Performs z-score normalization if isNormalized is 1
106     # Print intermidate training loss if verbose = True
107     def __init__(self,lr=0.005, num_iter=10000, verbose = True):
108         self.lr = lr
109         self.num_iter = num_iter
110         self.verbose = verbose
111         self.train_x = pd.DataFrame()
112         self.train_y = pd.DataFrame()
113         self.test_x = pd.DataFrame()
114         self.test_y = pd.DataFrame()
115         self.algType = 0
116         self.isNormalized = 0
117
118
119     def load_data(self, train_file, test_file):

```

```
120     self.train_x, self.train_y = getDataframe(train_file)
121     self.test_x, self.test_y = getDataframe(test_file)
122
123     def normalize(self):
124         # Applies z-score normalization to the dataframe and returns a normalized
dataframe
125         self.isNormalized = 1
126         data = np.append(self.train_x, self.test_x, axis = 0)
127         means = data.mean(0)
128         std = data.std(0)
129         self.train_x = (self.train_x - means).div(std)
130         self.test_x = (self.test_x - means).div(std)
131
132         # Gets the beta according to input
133     def train(self, algType):
134         self.algType = algType
135         newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-one column
as the first column
136         if(algType == '0'):
137             beta = getBeta_BatchGradient(newTrain_x, self.train_y.values, self.lr,
self.num_iter, self.verbose)
138             #print('Beta: ', beta)
139
140         elif(algType == '1'):
141             beta = getBeta_Newton(newTrain_x, self.train_y.values, self.num_iter,
self.verbose)
142             #print('Beta: ', beta)
143         else:
144             print('Incorrect beta_type! Usage: 0 - batch gradient descent, 1 -
Newton-Raphson method')
145
146         train_avglogL = compute_avglogL(newTrain_x, self.train_y.values, beta)
147         print('Training avgLogL: ', train_avglogL)
148
149         return beta
150
151     # Predict on given data x with learned parameter beta
152     def predict(self, x, beta):
153         newTest_x = addAllOneColumn(x)
154         self.predicted_y = (sigmoid(newTest_x.dot(beta))>=0.5)
155         return self.predicted_y
156
157     # predicted_y and y are the predicted and actual y values respectively as numpy
arrays
158     # function returns the accuracy
159     def compute_accuracy(self, predicted_y, y):
160         acc = np.sum(predicted_y == y)/predicted_y.shape[0]
161         return acc
162
163
```