

CS145 Howework 2

****Important Note:**** HW2 is due on **11:59 PM PT, Oct 30 (Friday, Week 4)**. Please submit through GradeScope.

Print Out Your Name and UID

****Name:** Kevin Li, **UID:** 405200619******

Before You Start

You need to first create HW2 conda environment by the given `cs145hw2.yml` file, which provides the name and necessary packages for this tasks. If you have `conda` properly installed, you may create, activate or deactivate by the following commands:

```
conda env create -f cs145hw2.yml
conda activate hw1
conda deactivate
```

OR

```
conda env create --name NAMEOFOURCHOICE -f cs145hw2.yml
conda activate NAMEOFOURCHOICE
conda deactivate
```

To view the list of your environments, use the following command:

```
conda env list
```

More useful information about managing environments can be found [here](https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html) (<https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html>).

You may also quickly review the usage of basic Python and Numpy package, if needed in coding for matrix operations.

In this notebook, you must not delete any code cells in this notebook. If you change any code outside the blocks (such as some important hyperparameters) that you are allowed to edit (between START/END YOUR CODE HERE), you need to highlight these changes. You may add some additional cells to help explain your results and observations.

```
In [12]: import numpy as np
import pandas as pd
import seaborn as sns
import sys
import random as rd
import matplotlib.pyplot as plt
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:
`%reload_ext autoreload`

If you can successfully run the code above, there will be no problem for environment setting.

1. Decision trees

This workbook will walk you through a decision tree.

1.1 Attribute selection measures

For classification models, misclassification rate is usually used as the final performance measurement. However, for classification trees, when selecting which attribute to split, measurements people often use includes information gain, gain ratio, and Gini index. Let's investigate these different measurements through the following problem.

Note: below shows how to calculate the misclassification rate of a classification tree with N total data points, K classes of the value we want to predict, and M leaf nodes.

In a node m , $m = 1, \dots, M$, let's denote the number of data points using N_m , and the number of data points in class k as N_{mk} , so the class prediction under majority vote is $j = \operatorname{argmax}_k N_{mk}$. The misclassification rate of this node m is $R_m = 1 - \frac{N_{mj}}{N_m}$. The total misclassification rate of the tree will be $R = \frac{\sum_{m=1}^M R_m * N_m}{N}$

Questions

Note: this question is a pure "question answer" problem. You don't need to do any coding.

Suppose our dataset includes a total of 800 people with 400 males and 400 females, and our goal is to do gender classification. Consider two different possible attributes we can split on in a decision tree model. Split on the first attribute results in a node11 with 300 male and 100 female, and a node12 with 100 male and 300 female. Split on the second attribute results in in a node21 with 400 male and 200 female, and a node22 with 200 female only.

1. Which split do you prefer when the measurement is misclassification rate and why?
2. What is the entropy in each of these four node?
3. What is the information gain of each of the two splits?
4. Which split do you prefer if the measurement is information gain. Do you see why it is an uncertainty or impurity measurement?
5. What is the gain ratio (normalized information gain) of each of the two splits? Which split do you prefer under this measurement. Do you get the same conclusion as information gain?

Your answer here:

Note: you can use several code cells to help you compute the results and answer the questions. Again you don't need to do any coding.

Please type your answer here!

1

The first split will classify all in node11 as male and node 12 as female yielding a misclassification rate of 1/4 The second split will classify all in node 21 as male and all in node22 as female, yielding a misclassification rate of 1/4 as well

Neither is superior

2

		A	B	C	D
1			Male	Female	Entropy
2	node11	300	100		0.81
3	node12	100	300		0.81
4	node21	400	200		0.92
5	node22	0	200		0

3

Original entropy: 1

Split 1:

$$1 - (400/800 \text{ Entropy}(\text{node11}) + 400/800 \text{ Entropy}(\text{node12}))$$

$$1 - 0.81$$

$$0.19$$

Split 2:

$$1 - (600/800 * \text{Entropy}(\text{node21}) + 200/800 * \text{Entropy}(\text{node22}))$$

$$1 - 0.69$$

$$0.31$$

4

The second split. This demonstrates how information gain is an impurity measurement since even with the same misclassification rate, the second is preferred due to it making a completely "pure" node (i.e. node22 is ONLY female)

5

Split 1:

InfoGain = 0.19

SplitInfo = $-(400/800 * \log_2(400/800) + 400/800 * \log_2(400/800))$

GainRatio = 0.19

Split 2:

InfoGain = 0.31

SplitInfo = $-(600/800 * \log_2(600/800) + 200/800 * \log_2(200/800))$

GainRatio = 0.38

Either works, but using gain ratio might not be necessary. This is because it tries to adjust for creating many branches/nodes due to a highly multivalued attribute. In our case, we don't have many nodes (only two nodes for either split). Gain ratio gives the same conclusion as information gain - this is because SplitInfo likes making fewer nodes and therefore likes things like large single nodes.

1.2 Coding decision trees

In this section, we are going to use the decision tree model to predict the the animal type class of the zoo dataset. The dataset has been preprocessed and splited into decision-tree-train.csv and decision-tree-test.csv for you.

```
In [13]: from hw2code.decision_tree import DecisionTree
mytree = DecisionTree()
mytree.load_data('./data/decision-tree-train.csv', './data/decision-tree-test.csv')
# As a sanity check, we print out the size of the training data (80, 17) and testing data (21, 17)
print('Training data shape: ', mytree.train_data.shape)
print('Testing data shape: ', mytree.test_data.shape)
```

Training data shape: (80, 17)

Testing data shape: (21, 17)

1.2.1 Information gain

Complete the `make_tree` and `compute_info_gain` function in `decision_tree.py`.

Train your model using `info_gain` measure to classify `type` and print the test accuracy.

```
In [14]: mytree = DecisionTree()
mytree.load_data('./data/decision-tree-train.csv', './data/decision-tree-test.csv')
test_acc = 0
#=====#
# START YOUR CODE HERE #
#=====#
mytree.train('type', 'info_gain')
test_acc = mytree.test('type')
#=====#
# END YOUR CODE HERE #
#=====#
print('Test accuracy is: ', test_acc)

best_feature is: legs
best_feature is: fins
best_feature is: toothed
best_feature is: eggs
best_feature is: hair
best_feature is: hair
best_feature is: toothed
best_feature is: aquatic
Test accuracy is: 0.8571428571428571
```

1.2.2 Gain ratio

Complete the `compute_gain_ratio` function in `decision_tree.py`.

Train your model using `gain_ratio` measure to classify `type` and print the test accuracy.

```
In [15]: mytree = DecisionTree()
mytree.load_data('./data/decision-tree-train.csv', './data/decision-tree-test.csv')
test_acc = 0
#####
# START YOUR CODE HERE #
#####
mytree.train('type', 'gain_ratio')
test_acc = mytree.test('type')
#####
# END YOUR CODE HERE #
#####
print('Test accuracy is: ', test_acc)

best_feature is: feathers
best_feature is: backbone
best_feature is: airborne
best_feature is: predator
best_feature is: milk
best_feature is: fins
best_feature is: legs
Test accuracy is: 0.8095238095238095
```

Question

Which measure do you like the most and why?

Your answer here:

I like information gain since it got a higher test accuracy.

2. SVM

This workbook will walk you through a SVM.

2.1 Support vectors and decision boundary

Note: for this question you can work entirely in the Jupyter Notebook, no need to edit any .py files.

Consider classifying the following 20 data points in the 2-d plane with class label y

```
In [16]: ds = pd.read_csv('data/svm-2d-data.csv')
ds.head()
# This command above will print out the first five data points
# in the dataset with column names as "x1", "x2" and "y"
# You may use command "ds" to show the entire dataset, which contains 20 data points
```

Out[16]:

	x1	x2	y
0	0.52	-1.00	1
1	0.91	0.32	1
2	-1.48	1.23	1
3	0.01	1.44	1
4	-0.46	-0.37	1

Suppose by solving the dual form of the quadratic programming of svm, we can derive the α_i 's for each data point as follows: Among $j = 0, 1, \dots, 19$ (note that the index starts from 0), $\alpha_1 = 0.5084$, $\alpha_5 = 0.4625$, $\alpha_{17} = 0.9709$, and $\alpha_j = 0$ for all other j .

Questions

1. Which vectors in the training points are support vectors?
2. What is the normal vector of the hyperplane w ?
3. What is the bias b ?
4. With the parameters w and b , we can now use our SVM to do predictions. What is predicted label of $x_{new} = (2, -0.5)$? Write out your $f(x_{new})$.
5. A plot of the data points has been generated for you. Please change the `support_vec` variable such that only the support vectors are indicated by red circles. Please also fill in the code to draw the decision boundary. Does your prediction of part 4 seems right visually on the plot?

Your answer here

Note: you can use several code cells to help you compute the results and answer the questions. Again you don't need to edit any .py files.

1

Points 1, 5, and 17 since they have non-zero alpha values

2

The normal for the hyperplane is [-1.34, -0.39]

```
In [17]: svs = ds.iloc[[1, 5, 17]]
vector = svs.drop(columns=['y']).multiply(svs['y'], axis='rows').multiply([.5084, .4625, .9709], axis='rows').sum()
vector
```

```
Out[17]: x1    -1.338076
          x2    -0.388998
          dtype: float64
```

3

With these support vectors, we know the distance to the hyperplane from each point is equivalent. Thus, we can turn those into a system of equations and solve it.

Distance from point to line where line is defined by a, b, β and point is defined by x, y is the following:

$$d = \frac{|ax+by+\beta|}{\sqrt{(a^2+b^2)}}$$

We know the a and b values since they're the normal vector, and by using one point with class=1 and one point with class=-1, we get the following:

$$d = \frac{|-1.338076 \cdot .91 - 0.388998 \cdot .32 + \beta|}{\sqrt{(1.338076^2 + 0.388998^2)}}$$

$$d = \frac{|-1.338076 \cdot 2.05 - 0.388998 \cdot 1.54 + \beta|}{\sqrt{(1.338076^2 + 0.388998^2)}}$$

We set them equal and solve

$$0.717631|\beta - 1.34213| = 0.717631|\beta - 3.34211|$$

Thus, we get the bias term (given that the normal is [-1.34, -0.39])

$$\beta = 2.34212$$

```
In [18]: for row in svs.iterrows():
          print('Point {}: ({}), ({}), class={}'.format(row[0], *row[1]))
```

```
Point 1: (0.91, 0.32), class=1.0
Point 5: (0.41, 2.04), class=1.0
Point 17: (2.05, 1.54), class=-1.0
```

4

$$f([2, -0.5]) = [2, -0.5, 1]^T [-1.34, -0.389, 2.34]$$

$$f([2, -0.5]) = -0.1455$$

The SVM predicts class=-1


```

In [19]: # answer 5
x1_range = np.arange(-2, 5, 0.5)
x2_range = np.arange(-2, 4., 0.5)

fig, ax = plt.subplots(figsize=(10, 8))
ax = fig.gca()
ax.set_xticks(x1_range)
ax.set_yticks(x2_range)
ax.grid()
ax.scatter(ds['x1'], ds['x2'], c=ds['y'])

support_vec = ds.iloc[[1, 5, 17]]
#####
# START YOUR CODE HERE #
#####
pt = (1.5, 0.861202)
ax.axline(pt, slope=-3.4398)
#####
# END YOUR CODE HERE #
#####
ax.scatter(support_vec['x1'], support_vec['x2'], marker='o', facecolor='none',
s=200, color='red')
sns.despine(ax=ax, left=True, bottom=True, offset=0)
plt.show()

```



5

The prediction looks right, and just barely as the numbers implied.

2.2 Coding SVM

In this section, we are going to use SVM for classifying the y value of 4-dimensional data points. The dataset has been preprocessed and splited into `svm-train.csv` and `svm-test.csv` for you.

For this question we are going to use the `cvxopt` package to help us solve the optimization problem of SVM. You will see it in the `.py` files, but you don't need to any coding with it. For this question, you only need to implement the right kernel function, and your kernel matrix `K` in `svm.py` line 135 will be plugged in the `cvxopt` optimization problem solver.

For more information about `cvxopt` please refer to <http://cvxopt.org/> (<http://cvxopt.org/>)

```
In [20]: from hw2code.svm import SVM
          svm = SVM()
          svm.load_data('./data/svm-train.csv', './data/svm-test.csv')
          # As a sanity check, we print out the size of the training data (1098, 4) and
          # (1098,) and testing data (274, 4) and (274,)
          print('Training data shape: ', svm.train_x.shape, svm.train_y.shape)
          print('Testing data shape:', svm.test_x.shape, svm.test_y.shape)
```

```
Training data shape: (1098, 4) (1098,)
Testing data shape: (274, 4) (274,)
```

2.2.1 Linear kernel

Complete the `SVM.predict` and `linear_kernel` function in `svm.py`. Train a hard margin SVM and a soft margin SVM with linear kernel. Print the test accuracy for both cases.

```
In [21]: svm_hard = SVM()
svm_hard.load_data('./data/svm-train.csv', './data/svm-test.csv')
hard_test_acc = 0
#####
# STRART YOUR CODE HERE #
#####
svm_hard.train()
hard_test_acc = svm_hard.test()
#####
# END YOUR CODE HERE #
#####

svm_soft = SVM()
svm_soft.load_data('./data/svm-train.csv', './data/svm-test.csv')
soft_test_acc = 0
#####
# STRART YOUR CODE HERE #
#####
svm_soft.train(C=1)
soft_test_acc = svm_soft.test()
#####
# END YOUR CODE HERE #
#####
print('Hard margin test accuracy is: ', hard_test_acc)
print('Soft margin test accuracy is: ', soft_test_acc)
```

```
1098 support vectors out of 1098 points
38 support vectors out of 1098 points
Hard margin test accuracy is: 0.5547445255474452
Soft margin test accuracy is: 0.9890510948905109
```

Questions

Are these two results similar? Why or why not?

Your Answer

The results are very different. This is because the hard SVM couldn't give any slack and was unable to split the dataset (i.e. the dataset is not linearly separable); it seems it ended up using every single point as a support vector. It likely failed to converge on the optimization problem successfully and correctly.

Because of this, it failed to intelligently pick a bounding hyperplane that was relevant to the data i.e. it did not find a hyperplane that intelligently split the dataset.

2.2.2 Polynomial kernel

Complete the `polynomial_kernel` function in `svm.py`. Train a soft margin SVM with degree 3 polynomial kernel and parameter `C = 100` for the regularization term. Print the test accuracy.

```
In [43]: svm = SVM()
svm.load_data('./data/svm-train.csv', './data/svm-test.csv')
test_acc = 0
#=====#
# START YOUR CODE HERE #
#=====#
svm.train(kernel_name='polynomial_kernel', C=100)
test_acc = svm.test()
#=====#
# END YOUR CODE HERE #
#=====#
print('Test accuracy is: ', test_acc)
```

```
19 support vectors out of 1098 points
Test accuracy is: 0.927007299270073
```

Questions

Is the result better than linear kernel? Why or why not?

Your Answer

It ended up doing slightly worse - this is probably because it overfit to the data as a polynomial kernel of is very similar to a linear kernel except that the higher the degree, the more closely it can fit the training data.

Since it was able to fit the training data better, it obviously chose to do that but ended it ended up slightly overfitting to the training data.

2.2.3 Gaussian kernel

Complete the `gaussian_kernel` function using the `gaussian_kernel_point` in `svm.py`. Train a soft margin SVM with Gaussian kernel and parameter `C = 100` for the regularization term. Print the test accuracy.

```
In [46]: svm = SVM()
svm.load_data('./data/svm-train.csv', './data/svm-test.csv')
test_acc = 0
#=====#
# STRART YOUR CODE HERE #
#=====#
svm.train(kernel_name='gaussian_kernel', C=100)
test_acc = svm.test()
#=====#
# END YOUR CODE HERE #
#=====#
print('Test accuracy is: ', test_acc)
```

35 support vectors out of 1098 points
Test accuracy is: 1.0

Questions

1. Is the result better than linear kernel and polynomial kernel? Why or why not?
2. Which one of these four models do you like the most and why?
3. (Bonus question, optional) Can you come up with a vectorized implementation of `gaussian_kernel` without calling `gaussian_kernel_point` ? Fill that in `svm.py`.

Your Answer

1

The gaussian kernel ended up being the most successful kernel - gaussian RBF kernels are far more powerful than linear kernels and allow separation of non-linearly separable data but they tend to not overfit as grossly as high degree polynomial kernels.

2

Gaussian RBF. It generalizes extremely well to real-world data, is smooth and doesn't overfit, and has been popular for a long time for those reasons.

3

Question 3 was done as can be seen in the `svm.py` file

End of Homework 2 :)

After you've finished the homework, please print out the entire `ipynb` notebook and two `py` files into one PDF file. Make sure you include the output of code cells and answers for questions. Prepare submit it to GradeScope. Also this time remember assign the pages to the questions on GradeScope


```

1 import pandas as pd
2 import numpy as np
3 from pprint import pprint
4 import sys
5
6 # Reads the data from CSV files, each attribute column can be obtained via its name,
  e.g., y = data['y']
7 def getDataframe(filePath):
8     data = pd.read_csv(filePath)
9     return data
10
11 # predicted_y and y are the predicted and actual y values respectively as numpy
  arrays
12 # function prints the accuracy
13 def compute_accuracy(predicted_y, y):
14     acc = 100.0
15     acc = np.sum(predicted_y == y)/predicted_y.shape[0]
16     return acc
17
18 #Compute entropy according to y distribution
19 def compute_entropy(y):
20     entropy = 0.0
21     elements, counts = np.unique(y, return_counts = True)
22     n = y.shape[0]
23
24     for i in range(len(elements)):
25         prob = counts[i]/n
26         if prob!= 0:
27             entropy -= prob * np.log2(prob)
28     return entropy
29
30 #att_name: attribute name; y_name: the target attribute name for classification
31 def compute_info_gain(data, att_name, y_name):
32     info_gain = 0.0
33
34     #Calculate the values and the corresponding counts for the select attribute
35     vals, counts = np.unique(data[att_name], return_counts=True)
36     total_counts = np.sum(counts)
37
38     #Calculate the conditional entropy
39     #=====#
40     # STRART YOUR CODE HERE #
41     #=====#
42     info_gain = compute_entropy(data[y_name])
43     info_gain -= sum(map(lambda x: x[1] / total_counts *
  compute_entropy(data.loc[data[att_name] == x[0]][y_name]), zip(vals, counts)))
44     #=====#
45     #   END YOUR CODE HERE   #
46     #=====#
47
48     return info_gain
49
50
51 def comput_gain_ratio(data, att_name, y_name):
52     gain_ratio = 0.0
53     #Calculate the values and the corresponding counts for the select attribute
54     vals, counts = np.unique(data[att_name], return_counts=True)
55     total_counts = np.sum(counts)
56
57     #Calculate the information for the selected attribute

```

```

58     att_info = 0.0
59     #=====#
60     # STRART YOUR CODE HERE #
61     #=====#
62     att_info -= sum(map(lambda count: count/total_counts *
np.log2(count/total_counts), counts))
63     #=====#
64     #   END YOUR CODE HERE   #
65     #=====#
66     gain_ratio = 0.0 if np.abs(att_info) < 1e-9 else min(1, compute_info_gain(data,
att_name, y_name) / att_info)
67     return gain_ratio
68
69 # Class of the decision tree model based on the ID3 algorithm
70 class DecisionTree(object):
71     def __init__(self):
72         self.train_data = pd.DataFrame()
73         self.test_data = pd.DataFrame()
74
75     def load_data(self, train_file, test_file):
76         self.train_data = getDataframe(train_file)
77         self.test_data = getDataframe(test_file)
78
79     def train(self, y_name, measure, parent_node_class= None):
80         self.y_name = y_name
81         self.measure = measure
82         self.tree = self.make_tree(self.train_data, parent_node_class)
83
84     def make_tree(self, train_data, parent_node_class = None):
85         data = train_data
86         features = data.drop(self.y_name, axis = 1).columns.values
87         measure = self.measure
88         #Stopping condition 1: If all target_values have the same value, return this
value
89         if len(np.unique(data[self.y_name])) <= 1:
90             leaf_value = -1
91             #=====#
92             # STRART YOUR CODE HERE #
93             #=====#
94             if len(np.unique(data[self.y_name])) > 0:
95                 leaf_value = np.unique(data[self.y_name])[0]
96             #=====#
97             #   END YOUR CODE HERE   #
98             #=====#
99             return leaf_value
100
101         #Stopping condition 2: If the dataset is empty, return the parent_node_class
102         elif len(data)== 0:
103             return parent_node_class
104
105         #Stopping condition 3: If the feature space is empty, return the majority
class
106         elif len(features) == 0:
107             return np.unique(data[self.y_name])
[ np.argmax(np.unique(data[y_name], return_counts=True)[1])]
108
109         # Not a leaf node, create an internal node
110         else:
111             #Set the default value for this node --> The mode target feature value
of the current node

```

```

112     parent_node_class = np.unique(data[self.y_name])
113     [np.argmax(np.unique(data[self.y_name], return_counts=True)[1])]
114     #Select the feature which best splits the dataset
115     if measure == 'info_gain':
116         item_values = [compute_info_gain(data, feature, self.y_name) for
feature in features] #Return the information gain values for the features in the
dataset
117     elif measure == 'gain_ratio':
118         item_values = [comput_gain_ratio(data, feature, self.y_name) for
feature in features] #Return the gain_ratio for the features in the dataset
119     else:
120         raise ValueError("kernel not recognized")
121
122     best_feature_index = np.argmax(item_values)
123     best_feature = features[best_feature_index]
124     print('best_feature is: ', best_feature)
125
126     #Create the tree structure. The root gets the name of the feature
(best_feature)
127     tree = {best_feature: {}}
128
129
130     #Grow a branch under the root node for each possible value of the root node
feature
131
132     for value in np.unique(data[best_feature]):
133         #Split the dataset along the value of the feature with the largest
information gain and therwith create sub_datasets
134         sub_data = data.where(data[best_feature] == value).dropna()
135
136         #Remove the selected feature from the feature space
137         sub_data = sub_data.drop(best_feature, axis = 1)
138
139         #Call the ID3 algorithm for each of those sub_datasets with the new
parameters --> Here the recursion comes in!
140         subtree = self.make_tree(sub_data, parent_node_class)
141
142         #Add the sub tree, grown from the sub_dataset to the tree under the root
node
143         tree[best_feature][value] = subtree
144
145     return tree
146
147
148     def test(self, y_name):
149         accuracy = self.classify(self.test_data, y_name)
150         return accuracy
151
152     def classify(self, test_data, y_name):
153         #Create new query instances by simply removing the target feature column
from the test dataset and
154         #convert it to a dictionary
155         test_x = test_data.drop(y_name, axis=1)
156         test_y = test_data[y_name]
157
158         n = test_data.shape[0]
159         predicted_y = np.zeros(n)
160
161         #Calculate the prediction accuracy

```

```
162         for i in range(n):
163             predicted_y[i] = DecisionTree.predict(self.tree, test_x.iloc[i])
164
165         output = np.zeros((n,2))
166         output[:,0] = test_y
167         output[:,1] = predicted_y
168         accuracy = compute_accuracy(predicted_y, test_y.values)
169         return accuracy
170
171     def predict(tree, query):
172         # find the root attribute
173         default = -1
174         for root_name in list(tree.keys()):
175             try:
176                 subtree = tree[root_name][query[root_name]]
177             except:
178                 return default ## root_name does not appear in query attribute list
179             (it is an error!)
180
181             ##if subtree is still a dictionary, recursively test next attribute
182             if isinstance(subtree,dict):
183                 return DecisionTree.predict(subtree, query)
184             else:
185                 leaf = subtree
186                 return leaf
187
```

```

1 import numpy as np
2 from numpy import linalg
3 import cvxopt
4 import cvxopt.solvers
5 import sys
6 import pandas as pd
7 cvxopt.solvers.options['show_progress'] = False
8
9 # Reads the data from CSV files, converts it into Dataframe and returns x and y
  dataframes
10 def getDataframe(filePath):
11     dataframe = pd.read_csv(filePath)
12     y = dataframe['y']
13     x = dataframe.drop('y', axis=1)
14     y = y*2 -1.0
15     return x.to_numpy(), y.to_numpy()
16
17 def compute_accuracy(predicted_y, y):
18     acc = 100.0
19     acc = np.sum(predicted_y == y)/predicted_y.shape[0]
20     return acc
21
22 def gaussian_kernel_point(x, y, sigma=5.0):
23     return np.exp(-linalg.norm(x-y)**2 / (2 * (sigma ** 2)))
24
25 def linear_kernel(X, Y=None):
26     Y = X if Y is None else Y
27     m = X.shape[0]
28     n = Y.shape[0]
29     assert X.shape[1] == Y.shape[1]
30     kernel_matrix = np.zeros((m, n))
31     #=====#
32     # STRART YOUR CODE HERE #
33     #=====#
34     kernel_matrix = X @ Y.T
35     #=====#
36     #   END YOUR CODE HERE   #
37     #=====#
38     return kernel_matrix
39
40 def polynomial_kernel(X, Y=None, degree=3):
41     Y = X if Y is None else Y
42     m = X.shape[0]
43     n = Y.shape[0]
44     assert X.shape[1] == Y.shape[1]
45     kernel_matrix = np.zeros((m, n))
46     #=====#
47     # STRART YOUR CODE HERE #
48     #=====#
49     kernel_matrix = ((X @ Y.T) + 1) ** degree
50     #=====#
51     #   END YOUR CODE HERE   #
52     #=====#
53     return kernel_matrix
54
55 # def gaussian_kernel(X, Y=None, sigma=5.0):
56 #     Y = X if Y is None else Y
57 #     m = X.shape[0]
58 #     n = Y.shape[0]
59 #     assert X.shape[1] == Y.shape[1]

```

```

60 #     kernel_matrix = np.zeros((m, n))
61 #     #=====#
62 #     # STRART YOUR CODE HERE #
63 #     #=====#
64 #     #=====#
65 #     #   END YOUR CODE HERE   #
66 #     #=====#
67 #     return kernel_matrix
68
69
70 # Bonus question: vectorized implementation of Gaussian kernel
71 # If you decide to do the bonus question, comment the gaussian_kernel function
  above,
72 # then implement and uncomment this one.
73 def gaussian_kernel(X, Y=None, sigma=5.0):
74     Y = X if Y is None else Y
75     assert X.shape[1] == Y.shape[1]
76
77     dists = np.linalg.norm(X[:,None,:] - Y[None,:,:], axis=-1)
78     kernel_matrix = np.exp(-dists**2/(2 * sigma**2))
79
80     return kernel_matrix
81
82 class SVM(object):
83     def __init__(self):
84         self.train_x = pd.DataFrame()
85         self.train_y = pd.DataFrame()
86         self.test_x = pd.DataFrame()
87         self.test_y = pd.DataFrame()
88         self.kernel_name = None
89         self.kernel = None
90
91     def load_data(self, train_file, test_file):
92         self.train_x, self.train_y = getDataframe(train_file)
93         self.test_x, self.test_y = getDataframe(test_file)
94
95
96     def train(self, kernel_name='linear_kernel', C=None):
97         self.kernel_name = kernel_name
98         if(kernel_name == 'linear_kernel'):
99             self.kernel = linear_kernel
100         elif(kernel_name == 'polynomial_kernel'):
101             self.kernel = polynomial_kernel
102         elif(kernel_name == 'gaussian_kernel'):
103             self.kernel = gaussian_kernel
104         else:
105             raise ValueError("kernel not recognized")
106
107         self.C = C
108         if self.C is not None:
109             self.C = float(self.C)
110
111         self.fit(self.train_x, self.train_y)
112
113     # predict labels for test dataset
114     def predict(self, X):
115         if self.w is not None: ## linear case
116             n = X.shape[0]
117             predicted_y = np.zeros(n)
118             #=====#

```

```

119         # STRART YOUR CODE HERE #
120         #=====#
121         predicted_y = self.kernel(X, self.w[:, None].T).T[0] + self.b
122         #=====#
123         # END YOUR CODE HERE #
124         #=====#
125         return predicted_y
126
127     else: ## non-linear case
128         n = X.shape[0]
129         predicted_y = np.zeros(n)
130         #=====#
131         # STRART YOUR CODE HERE #
132         #=====#
133         predicted_y = (self.a * self.sv_y * self.kernel(X, self.sv)).sum(axis=1)
+ self.b
134         #=====#
135         # END YOUR CODE HERE #
136         #=====#
137         return predicted_y
138
139         #=====#
140         # Please DON'T change any code below this line! #
141         #=====#
142     def fit(self, X, y):
143         n_samples, n_features = X.shape
144         # Kernel matrix
145         K = self.kernel(X)
146
147         # dealing with dual form quadratic optimization
148         P = cvxopt.matrix(np.outer(y,y) * K)
149         q = cvxopt.matrix(np.ones(n_samples) * -1)
150         A = cvxopt.matrix(y, (1,n_samples),'d')
151         b = cvxopt.matrix(0.0)
152
153         if self.C is None:
154             G = cvxopt.matrix(np.diag(np.ones(n_samples) * -1))
155             h = cvxopt.matrix(np.zeros(n_samples))
156         else:
157             tmp1 = np.diag(np.ones(n_samples) * -1)
158             tmp2 = np.identity(n_samples)
159             G = cvxopt.matrix(np.vstack((tmp1, tmp2)))
160             tmp1 = np.zeros(n_samples)
161             tmp2 = np.ones(n_samples) * self.C
162             h = cvxopt.matrix(np.hstack((tmp1, tmp2)))
163
164         # solve QP problem
165         solution = cvxopt.solvers.qp(P, q, G, h, A, b)
166         # Lagrange multipliers
167         a = np.ravel(solution['x'])
168
169         # Support vectors have non zero lagrange multipliers
170         sv = a > 1e-5
171         ind = np.arange(len(a))[sv]
172         self.a = a[sv]
173         self.sv = X[sv]
174         self.sv_y = y[sv]
175
176         print("%d support vectors out of %d points" % (len(self.a), n_samples))
177

```

```
178     # Intercept via average calculating b over support vectors
179     self.b = 0
180     for n in range(len(self.a)):
181         self.b += self.sv_y[n]
182         self.b -= np.sum(self.a * self.sv_y * K[ind[n],sv])
183     self.b /= len(self.a)
184
185     # Weight vector
186     if self.kernel_name == 'linear_kernel':
187         self.w = np.zeros(n_features)
188         for n in range(len(self.a)):
189             self.w += self.a[n] * self.sv_y[n] * self.sv[n]
190     else:
191         self.w = None
192
193
194     def test(self):
195         accuracy = self.classify(self.test_x, self.test_y)
196         return accuracy
197
198     def classify(self, X, y):
199         predicted_y = np.sign(self.predict(X))
200         accuracy = compute_accuracy(predicted_y, y)
201         return accuracy
```